

wrangle_act

September 25, 2019

```
In [1]: import datetime as dt
import json
import numpy as np
import pandas as pd
import re
import requests
import tweepy
```

Gather

```
In [2]: archive = pd.read_csv("twitter-archive-enhanced.csv")
archive.set_index("tweet_id", inplace = True)
archive.head(2)
```

```
Out[2]:
```

tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
892420643555336193	NaN	NaN	
892177421306343426	NaN	NaN	

tweet_id	timestamp	\
892420643555336193	2017-08-01 16:23:56 +0000	
892177421306343426	2017-08-01 00:17:27 +0000	

tweet_id	source	\
892420643555336193	<a href="http://twitter.com/download/iphone" r...	
892177421306343426	<a href="http://twitter.com/download/iphone" r...	

tweet_id	text	\
892420643555336193	This is Phineas. He's a mystical boy. Only eve...	
892177421306343426	This is Tilly. She's just checking pup on you...	

tweet_id	retweeted_status_id	retweeted_status_user_id	\
892420643555336193	NaN	NaN	
892177421306343426	NaN	NaN	

retweeted_status_timestamp \

```
expanded_urls \
```

```
rating_numerator rating_denominator name doggo \
```

floofer pupper puppo

```
r = requests.get(tsv_url)
```

```
with open(tsv_url.split('/')[-1], mode = 'wb') as file:
```

```
images = pd.read_csv('image-predictions.tsv', sep = '\t')
```

```
Out[3]:
```

0 666020888022790149 <https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg>

img_num	p1	p1_conf	p1_dog	p2	\
---------	----	---------	--------	----	---

p2_conf p2_dog p3 p3_conf p3_dog

```
In [4]: consumer_key = 'IudBZvlwcR2jN5zeo0wA90TvB'
```

```
auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
```

```

auth.set_access_token(access_token, access_token_secret)

api = tweepy.API(auth, parser=tweepy.parsers.JSONParser(), wait_on_rate_limit=True)

In [ ]: df = []
        exceptions = []
        tweet_id = images['tweet_id']

        for id in tweet_id:
            try:
                page = api.get_status(id)
                favorites = page['favorite_count']
                retweets = page['retweet_count']
                time = pd.to_datetime(page['created_at'])
                df.append({'tweet_id': int(id),
                           'favorites': int(favorites),
                           'retweets': int(retweets)})

            except Exception as e:
                exceptions.append(id)

In [ ]: exceptions

In [ ]: exceptions2 = []
        for e in exceptions:
            try:
                page = api.get_status(e)
                favorites = page['favorite_count']
                retweets = page['retweet_count']
                time = pd.to_datetime(page['created_at'])
                df.append({'tweet_id': int(e),
                           'favorites': int(favorites),
                           'retweets': int(retweets)})

            except Exception:
                exceptions2.append(id)

In [ ]: df = pd.DataFrame(df, columns = ['tweet_id', 'favorites', 'retweets'])
        df.to_csv('tweet_json.txt', encoding = 'utf-8')

In [ ]: df = pd.read_csv('tweet_json.txt', encoding = 'utf-8', index = False)
        df.set_index('tweet_id', inplace = True)
        df.tail()

In [ ]: images.set_index('tweet_id', inplace = True)
        df2 = pd.merge(left=archive, right=images, left_index=True, right_index=True, how='left')
        df2 = pd.merge(left=df2, right=df, left_index=True, right_index=True, how='left')
        df2.to_csv('df2copy.csv', encoding = 'utf-8')

```

```

In [ ]: #Assessing Data
        archive.info()

In [ ]: archive.name.value_counts()

In [ ]: archive.rating_denominator.value_counts()

In [ ]: archive.rating_numerator.value_counts()

In [ ]: archive.head()

In [ ]: images.info()

In [ ]: tweet = pd.read_csv("tweet_json.txt", encoding = 'utf-8', index = 'False')
        tweet.info()

```

Quality Several columns have empty values, like `in_reply_to_status`, `in_reply_to_user_id`, `retweeted_status_id`, `retweeted_status_user_id`, `retweeted_status_timestamp`. The `name` column has many entries which do not look like names. The most frequent entry in `name` column is "a", which is not a name. The `numerator` and `denominator` columns have unusual values. The `timestamp` column is an object. It has to be a datetime object. There are 2075 rows in the `images` dataframe and 2356 rows in the `archive` dataframe. In several columns, null values are not treated as null values. Tidiness The dog stages have values as columns, instead of one column filled with their values. We don't need the `Unnamed: 0` column from the `tweet` dataframe. The columns for dog breed predictions can be condensed.

```

In [ ]: df = pd.read_csv("df2copy.csv")

In [ ]: Clean

In [ ]: #delete extra column
        del(df['Unnamed: 0'])

In [ ]: #test
        df.columns

In [ ]: #convert timestamp to datetime
        df['timestamp'] = pd.to_datetime(df['timestamp'])

In [ ]: df.info()

In [ ]: # removing retweets
        df = df[pd.isnull(df['retweeted_status_id'])]
        df.shape[0]

In [ ]: df.drop(['retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp'])

In [ ]: #test
        df.columns

```

```

In [ ]: # Condensing dog type columns
dog_type = []

x = ['pupper', 'puppo', 'doggo', 'floof']
y = ['pupper', 'puppo', 'doggo', 'floof']

for row in df['text']:
    row = row.lower()
    for word in x:
        if word in str(row):
            dog_type.append(y[x.index(word)])
            break
    else:
        dog_type.append('None')

df['dog_type'] = dog_type

In [ ]: df['dog_type'].value_counts()

In [ ]: # removing extra columns
df.drop(['doggo', 'floofer', 'pupper', 'puppo'], axis=1, inplace=True)

In [ ]: #test
df.columns

In [ ]: #Condensing dog breed predictions
breed = []
conf= []

def breed_conf(row):
    if row['p1_dog']:
        breed.append(row['p1'])
        conf.append(row['p1_conf'])
    elif row['p2_dog']:
        breed.append(row['p2'])
        conf.append(row['p2_conf'])
    elif row['p3_dog']:
        breed.append(row['p3'])
        conf.append(row['p3_conf'])
    else:
        breed.append('Unidentifiable')
        conf.append(0)

df.apply(breed_conf, axis = 1)

df['breed'] = breed
df['confidence'] = conf

In [ ]: #removing the processed columns
df.drop(['p1', 'p1_conf', 'p1_dog', 'p2', 'p2_conf', 'p2_dog', 'p3', 'p3_conf', 'p3_dog']

```

```

In [ ]: #test
        df.head(2)

In [ ]: #removing useless columns
        df['in_reply_to_status_id'].value_counts()

In [ ]: df['in_reply_to_user_id'].value_counts()

In [ ]: df.drop(['in_reply_to_status_id', 'in_reply_to_user_id'], axis=1, inplace=True)

In [ ]: #test
        df.columns

In [ ]: # Extract Dog Rates and Dog Count
        rates = []

        #raw_rates = lambda x: rates.append(re.findall(r'(\d+(\.\d+))|(\d+)\|(\d+0)', x, flags=0))

        df['text'].apply(lambda x: rates.append(re.findall(r'(\d+(\.\d+))|(\d+)\|(\d+0)', x, flags=0)))

        rating = []
        dog_count = []

        for item in rates:

            # for tweets with no rating, but a picture, so a dog_count of 1
            if len(item) == 0:
                rating.append('NaN')
                dog_count.append(1)

            # for tweets with single rating and dog_count of 1
            elif len(item) == 1 and item[0][-1] == '10':
                rating.append(float(item[0][0]))
                dog_count.append(1)

            # for multiple ratings
            elif len(item) == 1:
                a = float(item[0][0]) / (float(item[0][-1]) / 10)
                rating.append(a)
                dog_count.append(float(item[0][-1]) / 10)
                # for tweets with more than one rating
            elif len(item) > 1:
                total = 0
                r = []
                for i in range(len(item)):
                    if item[i][-1] == '10': #one tweet has the phrase '50/50' so I'm coding to e
                        r.append(item[i])
                for rate in r:
                    total = total + float(rate[0])

```

```

        a = total / len(item)
        rating.append(a)
        dog_count.append(len(item))

    # if any error has occurred
    else:
        rating.append('Not parsed')
        dog_count.append('Not parsed')

df['rating'] = rating # not need to also add denominator since they are all 10!
df['dog_count'] = dog_count
df['rating'].value_counts()

In [ ]: df.drop(['rating_numerator', 'rating_denominator'], axis=1, inplace=True)

In [ ]: #test
        df.info()

In [ ]: df['dog_count'].value_counts()

In [ ]: # extract names
        df['text_split'] = df['text'].str.split()

In [ ]: names = []

        # use string starts with method to clean this up

def extract_names(row):

    # 'named Phineas'
    if 'named' in row['text'] and re.match(r'[A-Z].*', row['text_split'][(row['text_split'].index('named') + 1)]):
        names.append(row['text_split'][(row['text_split'].index('named') + 1)]) # 'H
    elif row['text'].startswith('Here we have ') and re.match(r'[A-Z].*', row['text_split'][(row['text_split'].index('Here we have ') + 1)]):
        names.append(row['text_split'][(row['text_split'].index('Here we have ') + 1)].strip('.').strip(','))

    # 'This is Phineas'
    elif row['text'].startswith('This is ') and re.match(r'[A-Z].*', row['text_split'][(row['text_split'].index('This is ') + 1)]):
        names.append(row['text_split'][(row['text_split'].index('This is ') + 1)].strip('.').strip(','))

    # 'Say hello to Phineas'
    elif row['text'].startswith('Say hello to ') and re.match(r'[A-Z].*', row['text_split'][(row['text_split'].index('Say hello to ') + 1)]):
        names.append(row['text_split'][(row['text_split'].index('Say hello to ') + 1)].strip('.').strip(','))

    # 'Meet Phineas'
    elif row['text'].startswith('Meet ') and re.match(r'[A-Z].*', row['text_split'][(row['text_split'].index('Meet ') + 1)]):
        names.append(row['text_split'][(row['text_split'].index('Meet ') + 1)].strip('.').strip(','))

    else:
        names.append('Nameless')

```

```

df.apply(extract_names, axis=1)

df['names'] = names

In [ ]: df['names'].value_counts()

In [ ]: df.drop(['text_split'], axis=1, inplace=True)

In [ ]: df.loc[df['names'] == 'Nameless', 'names'] = None
df.loc[df['breed'] == 'Unidentifiable', 'breed'] = None
df.loc[df['dog_type'] == 'None', 'dog_type'] = None
df.loc[df['rating'] == 0.0, 'rating'] = np.nan
df.loc[df['confidence'] == 0.0, 'confidence'] = np.nan

In [ ]: #test
df.info()

In [ ]: #saving the cleaned file
df.to_csv('twitter_archive_master.csv', encoding = 'utf-8')

In [ ]: Analysis

In [ ]: %matplotlib inline
import matplotlib
import matplotlib.pyplot as plt

In [ ]: df = pd.read_csv('twitter_archive_master.csv')
df['timestamp'] = pd.to_datetime(df['timestamp'])
df.set_index('timestamp', inplace=True)

In [ ]: # Retweets, Favorites and Ratings Correlation
df[['favorites', 'retweets']].plot(style = '.', alpha = 0.4)
plt.title('Favorites and Retweets with Time')
plt.xlabel('Date')
plt.ylabel('Count');

In [ ]: df.plot(y='rating', ylim=[0,14], style = '.', alpha = 0.4)
plt.title('Rating with Time')
plt.xlabel('Date')
plt.ylabel('Rating');

In [ ]: df[['favorites', 'rating', 'retweets']].corr(method='pearson')

```

So Brant was right, there are more ratings above 10. Still don't know the reason why there are so much high ratings.

So let's see if dogs with higher ratings were getting more favorites and retweets. According to me, if the dogs are getting better they should be getting more favorites and retweets along with

the higher rating. There is a strong correlation between favorites and retweets. This means that if the tweet is good in general then there will be more retweets and favorites.

Yet there is no correlation between rating and retweets or rating and favorites. It can be because the dogs are not actually getting better. It can be that 'lower quality' dogs are given funnier captions. In this case, it is the caption that is getting more retweets and favorites, rather than the dog itself.

```
In [ ]: # Dog Stages Stats
        df.boxplot(column='rating', by='dog_type');

In [ ]: df.groupby('dog_type')['rating'].describe()

In [ ]: df.reset_index(inplace=True)
        df.groupby('dog_type')['timestamp'].describe()
```

So puppies are getting much lower rates than the other dog types. They have several low outliers which decrease the mean to 10.6.

Floofers are consistently rated above 10. I don't know whether they are really good or the rating just gets higher with time. Maybe we can see if 'floof' is a newer term.

Here we see that 'floof' is not a new term, first seen on January 2016. So we can say that floofer are consistently good dogs.

```
In [ ]: # Most Rated Breeds
        top=df.groupby('breed').filter(lambda x: len(x) >= 20)
        top['breed'].value_counts().plot(kind = 'bar')
        plt.title('The Most Rated Breeds');
```

It's difficult to know why these breeds are the top breeds. It could be because they are commonly owned. Or they could be the easiest to identify by the AI that identified them.

```
In [ ]: top.groupby('breed')['rating'].describe()

In [ ]: df['rating'].describe()

In [ ]: df[df['rating'] <= 14]['rating'].describe()
```

Here we have a statistical comparison of the top breeds with all the ratings. Only one of the top breeds has a mean higher than the total population mean. This is because of these two ratings: 420 and 1776.

Excluding outliers bring down the mean to 10.55.

```
In [ ]:
```